

Omnivore Ingestion Overview

Omnivore is a universal document ingestion pipeline. It accepts any file format, extracts structured intelligence from it, and stores the results in PostgreSQL so they can be searched semantically or queried like a database.

The extraction core normalizes every source into a neutral intermediate representation of blocks and fragments. A structure-first chunker then groups fragments along heading boundaries before embedding them with a local BGE model for hybrid retrieval.